Vaghesan Sundaram

(917) 297-6785 | vaghesan2@gmail.com | GitHub | LinkedIn | Portfolio

Education

University of Maryland, College Park **Expected Fall 2027**
B.S. Computer Science, Cybersecurity Focus; B.S. Mathematics; ACES Cybersecurity Minor College Park, MD
Relevant Coursework: Data Structures, Algorithms, Computer Systems, Statistics, Cryptography, Network Security, HACS408T Penetration Testing

Middlesex College **June 2024**
A.S. Mathematics Edison, NJ

Certification & Technical Skills

CompTIA Security+ (SY0-701), issued March 2026; expires March 2029.

Languages: Python, Java, C++
Security: Secure code review, vulnerability analysis, authentication review, evidence and severity triage, remediation reporting
Tools: Semgrep, Wireshark, Linux, Git, PowerShell; Burp Suite and Nmap in controlled labs

Experience

Cybersecurity Intern **Jun 2026 to Present**
Redmind Technologies *Remote*

- Assessed one public-facing client site and documented 7 findings spanning data exposure, authentication flows, configuration, and dependency risk.
- Combined Semgrep and manual review with evidence capture, severity triage, and remediation mapping to produce developer-facing guidance.
- Presented the findings in a written report and slide deck to developers, explaining risk, evidence, and recommended fixes.

Information Technology Intern **Jun to Aug 2025**
New Jersey Courts *Toms River, NJ*

- Supported production IT systems across court departments through troubleshooting, access workflows, user support, documentation, and escalation.
- Documented recurring issues and support procedures for repeatable troubleshooting and clearer team handoffs.

Selected Security Projects

Agent Least Privilege GitHub

- Built a 60-cell synthetic tool-use evaluation with gpt-5.4-mini; task-conditioned enforcement blocked all observed unauthorized state-changing effects (0/10) versus 4/10 unrestricted and 6/10 with a static allowlist.
- Found that enforcement prevented effects without eliminating model proposals: 5/10 monitored attack cases still contained an unauthorized proposed action.

CTI Claim Provenance GitHub

- Built a 192-cell gpt-5.6-luna comparison of citation-prompted and constrained evidence-grounding pipelines over frozen vulnerability-source questions.
- Measured a mixed result: evidence-binding improved from 83.3% to 86.5%, while strict canonical-answer matching declined from 33.3% to 26.0%.

Keystroke Authentication System GitHub

- Built a two-service Flask authentication system with SQLite persistence, typing-cadence validation, JWT-bound flows, bcrypt fallback, rate limiting, and single-use server-to-server code exchange.
- Added an end-to-end smoke test covering account creation, three-sample enrollment, and authentication.